

```

!pip install --upgrade pip
!pip install torch torchvision torchaudio
!pip install fsspec==2024.6.1
!pip install datasets==3.0.0
!pip install gcsfs==2024.6.0
!pip install jiwer
!pip install evaluate

# Imports
from google.colab import drive
import os, sys, itertools
import pandas as pd
from sklearn.model_selection import train_test_split
from PIL import Image

import torch
from torch.utils.data import Dataset

from datasets import load_dataset
import transformers
from transformers import Seq2SeqTrainingArguments, Seq2SeqTrainer
from transformers import VisionEncoderDecoderModel, TrOCRProcessor, default_data_collator
import evaluate

```

```

🔗 Requirement already satisfied: pip in /usr/local/lib/python3.10/dist-packages
Collecting pip
  Downloading pip-24.3.1-py3-none-any.whl.metadata (3.7 kB)
  Downloading pip-24.3.1-py3-none-any.whl (1.8 MB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.8/1.8 MB 55.0 MB/s eta 0:00:00
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 24.1.2
    Uninstalling pip-24.1.2:
      Successfully uninstalled pip-24.1.2
Successfully installed pip-24.3.1
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: torchaudio in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: mpmath<1.4, >=1.1.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pillow!=8.3.*, >=5.3.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages
Collecting fsspec==2024.6.1
  Downloading fsspec-2024.6.1-py3-none-any.whl.metadata (11 kB)
  Downloading fsspec-2024.6.1-py3-none-any.whl (177 kB)
Installing collected packages: fsspec

```

```

installing collected packages: fsspec
  Attempting uninstall: fsspec
    Found existing installation: fsspec 2024.10.0
    Uninstalling fsspec-2024.10.0:
      Successfully uninstalled fsspec-2024.10.0
ERROR: pip's dependency resolver does not currently take into account all the
gcsfs 2024.10.0 requires fsspec==2024.10.0, but you have fsspec 2024.6.1 which
Successfully installed fsspec-2024.6.1
Collecting datasets==3.0.0
  Downloading datasets-3.0.0-py3-none-any.whl.metadata (19 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/d
Collecting dill<0.3.9,>=0.3.0 (from datasets==3.0.0)
  Downloading dill-0.3.8-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.10/c
Requirement already satisfied: tqdm>=4.66.3 in /usr/local/lib/python3.10/dist-
Collecting xxhash (from datasets==3.0.0)
  Downloading xxhash-3.5.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86
Collecting multiprocessing (from datasets==3.0.0)
  Downloading multiprocessing-0.70.17-py310-none-any.whl.metadata (7.2 kB)
Requirement already satisfied: fsspec<=2024.6.1,>=2023.1.0 in /usr/local/lib/p
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: huggingface-hub>=0.22.0 in /usr/local/lib/pytho
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/pytho
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/c
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/pytho
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dis

```

```

# Environment info
print("Python:".rjust(15), sys.version[0:6])
print("Pandas:".rjust(15), pd.__version__)
print("Transformers:".rjust(15), transformers.__version__)
print("Torch:".rjust(15), torch.__version__)

```

```

Python: 3.10.1
Pandas: 2.2.2
Transformers: 4.46.3
Torch: 2.5.1+cu121

```

```

# Mount Google Drive
drive.mount('/content/drive', force_remount=True)
path = '/content/drive/My Drive/CMPE 252 Project/whiteplate_normal/'

```

```

Mounted at /content/drive

```

```

# Dataset Preparation
file_names, texts = [], []
for file in os.listdir(path):
    if file.endswith(('.jpg', '.png')):

```

```

if file.endswith( '.jpg' , '.png' ):
    file_names.append(file)
    texts.append(os.path.splitext(file)[0])

dataset = pd.DataFrame({'file_name': file_names, 'text': texts})
train_dataset, test_dataset = train_test_split(dataset, train_size=0.80, random_s
train_dataset.reset_index(drop=True, inplace=True)
test_dataset.reset_index(drop=True, inplace=True)

```

```

class License_Plates_OCR_Dataset(Dataset):
    def __init__(self, root_dir, df, processor, max_target_length=128):
        self.root_dir = root_dir
        self.df = df
        self.processor = processor
        self.max_target_length = max_target_length

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        file_name = self.df['file_name'][idx]
        text = self.df['text'][idx]
        image = Image.open(self.root_dir + file_name).convert("RGB")
        pixel_values = self.processor(image, return_tensors="pt").pixel_values
        labels = self.processor.tokenizer(text, padding="max_length", max_length=
        labels = [label if label != self.processor.tokenizer.pad_token_id else -1
        return {"pixel_values": pixel_values.squeeze(), "labels": torch.tensor(la

```

```
!huggingface-cli login
```

```

 _| _| _| _| _|_|_| _|_|_| _|_|_| _| _| _|_|_|
 _| _| _| _| _| _|_|_| _|_|_| _|_|_| _|_| _| _|_|_|
 _|_|_|_| _| _| _| _| _|_|_| _|_|_| _| _| _| _|_|_|
 _| _| _| _| _| _| _|_|_| _|_|_| _| _| _|_|_| _|_|_|
 _| _| _|_| _|_|_| _|_|_| _|_|_| _| _| _| _|_|_|

```

```

To log in, `huggingface_hub` requires a token generated from https://huggingface.co/settings/tokens
Enter your token (input will not be visible):
Add token as git credential? (Y/n) n
Token is valid (permission: fineGrained).
The token `testingLicensePlate` has been saved to /root/.cache/huggingface/sto
Your token has been saved to /root/.cache/huggingface/token
Login successful.
The current active token is: `testingLicensePlate`

```

```

# Model Initialization
MODEL_CKPT = "microsoft/trocr-base-printed"
processor = TrOCRProcessor.from_pretrained(MODEL_CKPT)
train_ds = License_Plates_OCR_Dataset(path, train_dataset, processor)

```

```
test_ds = License_Plates_UK_Dataset(path, test_dataset, processor)
```

```
model = VisionEncoderDecoderModel.from_pretrained(MODEL_CKPT)
model.config.decoder_start_token_id = processor.tokenizer.cls_token_id
model.config.pad_token_id = processor.tokenizer.pad_token_id
model.config.eos_token_id = processor.tokenizer.sep_token_id
model.config.max_length = 64
```

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access private repositories.
warnings.warn(

preprocessor_config.json: 100%  224/224 [00:00<00:00, 18.4kB/s]

tokenizer_config.json: 100%  1.12k/1.12k [00:00<00:00, 104kB/s]

vocab.json: 100%  899k/899k [00:00<00:00, 1.39MB/s]

merges.txt: 100%  456k/456k [00:00<00:00, 1.07MB/s]

special_tokens_map.json: 100%  772/772 [00:00<00:00, 67.0kB/s]

config.json: 100%  4.13k/4.13k [00:00<00:00, 334kB/s]

model.safetensors: 100%  1.33G/1.33G [00:05<00:00, 224MB/s]

```
Config of the encoder: <class 'transformers.models.vit.modeling_vit.ViTModel':
  "attention_probs_dropout_prob": 0.0,
  "encoder_stride": 16,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.0,
  "hidden_size": 768,
  "image_size": 384,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "model_type": "vit",
  "num_attention_heads": 12,
  "num_channels": 3,
  "num_hidden_layers": 12,
  "patch_size": 16,
  "qkv_bias": false,
  "transformers_version": "4.46.3"
}
```

```
Config of the decoder: <class 'transformers.models.trocr.modeling_trocr.TrOCRModel':
```

```

"activation_dropout": 0.0,
"activation_function": "gelu",
"add_cross_attention": true,
"attention_dropout": 0.0,
"bos_token_id": 0,
"classifier_dropout": 0.0,
"cross_attention_hidden_size": 768,
"d_model": 1024,
"decoder_attention_heads": 16,
"decoder_ffn_dim": 4096,
"decoder_layerdrop": 0.0,
"decoder_layers": 12,
"decoder_start_token_id": 2,
"dropout": 0.1,
"eos_token_id": 2,
"init_std": 0.02,
"is_decoder": true,
"layernorm_embedding": true,
"max_position_embeddings": 512,
"model_type": "trocr",
"pad_token_id": 1,
"scale_embedding": false,
"transformers_version": "4.46.3",
"use_cache": false,
"use_learned_position_embeddings": true,
"vocab_size": 50265
}

```

```

# Metrics
cer_metric = evaluate.load("cer")

# Adversarial Patch Attack to Maximize CER
def overlay_patch(image_tensor, patch):
    """Overlay patch on center of the image."""
    patched_image = image_tensor.clone()
    patch_height, patch_width = patch.shape[1:]
    center_y = (patched_image.shape[1] - patch_height) // 2
    center_x = (patched_image.shape[2] - patch_width) // 2
    patched_image[:, center_y:center_y + patch_height, center_x:center_x + patch_
    return patched_image

def fast_gradient_sign_patch(image_tensor, true_label, processor, model, cer metr
                                epochs=10, lr=0.01):
    """
    Generate an adversarial patch to fool the OCR model, ensuring device consiste

    # Ensure all computations happen on the same device
    device = model.device
    image_tensor = image_tensor.to(device)
    patch = torch.rand((3, *patch_size), device=device, requires_grad=True)
    optimizer = torch.optim.Adam([patch], lr=lr)

```

```

# Tokenize the true label and move it to the same device
true_ids = processor.tokenizer(true_label, return_tensors="pt").input_ids.to(device)

for epoch in range(epochs):
    optimizer.zero_grad()

    # Overlay the patch on the image
    patched_image = overlay_patch(image_tensor, patch)

    # Forward pass through the model
    outputs = model(patched_image.unsqueeze(0), labels=true_ids)

    # Use cross-entropy loss
    loss = outputs.loss
    loss.backward()
    optimizer.step()

    # Clamp patch values to ensure valid pixel range
    patch.data = torch.clamp(patch.data, 0, 1)

    # Compute CER (for logging purposes)
    predictions = processor.batch_decode(outputs.logits.argmax(dim=-1), skip_special_tokens=True)
    cer = cer_metric.compute(predictions=predictions, references=[true_label])

    print(f"Epoch {epoch + 1}/{epochs}, Loss: {loss.item():.4f}, CER: {cer:.4f}")

return patch

```

Downloading builder script: 100%  5.60k/5.60k [00:00<00:00, 456kB/s]

```

# Evaluate on Test Image
true_label = test_dataset.iloc[0]['text']
image_path = path + test_dataset.iloc[0]['file_name']

# Load and preprocess the image
image = Image.open(image_path).convert("RGB")
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Move processor-preprocessed tensor to the same device
image_tensor = processor(image, return_tensors="pt").pixel_values[0].to(device)

# Ensure the model is on the same device
model = model.to(device)

# Generate Adversarial Patch
adversarial_patch = fast_gradient_sign_patch(
    model, image_tensor, true_label, num_iter=100, num_restarts=10,
    random_restart_prob=0.1, verbose=True
)

```

```

    image_tensor, true_label, processor, model, cer_metric,
    patch_size=(100, 100), epochs=10, lr=0.01
)

# Apply the adversarial patch to the image
patched_image = overlay_patch(image_tensor, adversarial_patch)

# Ensure patched_image is on the correct device and has batch dimension
patched_image = patched_image.unsqueeze(0).to(device)

# Evaluate the model on the patched image
model.eval() # Set model to evaluation mode
with torch.no_grad(): # Disable gradient computation
    outputs = model.generate(patched_image) # Use generate for inference

# Decode predictions and compute CER
patched_predictions = processor.batch_decode(outputs, skip_special_tokens=True)
patched_cer = cer_metric.compute(predictions=patched_predictions, references=[true_label])

# Print results
print(f"True Label: {true_label}")
print(f"Predicted Label (Patched): {patched_predictions}")
print(f"CER on Patched Image: {patched_cer:.4f}")

```

```

Epoch 1/10, Loss: 6.1174, CER: 0.8571
Epoch 2/10, Loss: 5.8741, CER: 0.8571
Epoch 3/10, Loss: 5.7155, CER: 0.7143
Epoch 4/10, Loss: 5.6127, CER: 0.7143
Epoch 5/10, Loss: 5.5154, CER: 0.7143
Epoch 6/10, Loss: 5.4110, CER: 0.5714
Epoch 7/10, Loss: 5.3069, CER: 0.5714
Epoch 8/10, Loss: 5.2190, CER: 0.5714
Epoch 9/10, Loss: 5.1339, CER: 0.5714
Epoch 10/10, Loss: 5.0390, CER: 0.4286
/usr/local/lib/python3.10/dist-packages/transformers/generation/utils.py:1493:
  warnings.warn(
True Label: PZ68KRY
Predicted Label (Patched): ['@8 PZCSKRY']
CER on Patched Image: 0.7143

```

```

import matplotlib.pyplot as plt

def show_original_and_patched(original_image_tensor, patched_image_tensor, title1: str, title2: str):
    """
    Displaying the original and patched images side by side
    """
    # Remove batch dimension if it exists
    if len(patched_image_tensor.shape) == 4: # Shape [1, C, H, W]
        patched_image_tensor = patched_image_tensor.squeeze(0)
    # Tensor to NumPy array


```

```

#Tensors to Numpy arrays
original_image_np = original_image_tensor.permute(1, 2, 0).cpu().detach().numpy
patched_image_np = patched_image_tensor.permute(1, 2, 0).cpu().detach().numpy
#Plots
fig, axes = plt.subplots(1, 2, figsize=(12, 6))
#Original
axes[0].imshow(original_image_np)
axes[0].set_title(title1, fontsize=14)
axes[0].axis("off") #No axes
#Patched
axes[1].imshow(patched_image_np)
axes[1].set_title(title2, fontsize=14)
axes[1].axis("off") #No axes
plt.tight_layout()
plt.show()

```

show_original_and_patched(image_tensor, patched_image)



Api key: 610b65cafc807a2520e2754f9248364c728ef52b

```

# Compute metrics for training and evaluation
def compute_metrics(pred):
    label_ids = pred.label_ids
    pred_ids = pred.predictions

    # Decode predictions and labels
    pred_str = processor.batch_decode(pred_ids, skip_special_tokens=True)

```



```

    pred_str = processor.batch_decode(pred_ids, skip_special_tokens=True)
    label_ids[label_ids == -100] = processor.tokenizer.pad_token_id
    label_str = processor.batch_decode(label_ids, skip_special_tokens=True)

    cer = cer_metric.compute(predictions=pred_str, references=label_str)
    return {"cer": cer}

# Seq2Seq Training Arguments
MODEL_NAME = "trocr_license_plate_model"
NUM_OF_EPOCHS = 2

args = Seq2SeqTrainingArguments(
    output_dir=MODEL_NAME,
    num_train_epochs=NUM_OF_EPOCHS,
    predict_with_generate=True,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    per_device_train_batch_size=8,
    per_device_eval_batch_size=8,
    logging_first_step=True,
    hub_private_repo=True,
    push_to_hub=True
)

# Initialize Trainer
trainer = Seq2SeqTrainer(
    model=model,
    tokenizer=processor.feature_extractor,
    args=args,
    compute_metrics=compute_metrics,
    train_dataset=train_ds,
    eval_dataset=test_ds,
    data_collator=default_data_collator
)

# Train the model
print("\n-- Training the Model --")
train_results = trainer.train()
trainer.save_model()
trainer.log_metrics("train", train_results.metrics)
trainer.save_metrics("train", train_results.metrics)
trainer.save_state()

# Evaluate on Clean Test Set
print("\n-- Evaluating on Clean Test Set --")
metrics = trainer.evaluate()
trainer.log_metrics("eval", metrics)
trainer.save_metrics("eval", metrics)

```

```

/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1568: FutureWarning: warn(

```

